



Visual Odometry for Navigating in GPS Denied Environments

Final Report Summary

Alex Kreimer

Department of Computer Science

0.1 Introduction

In this work we revisit the problem of visual odometry. Visual odometry is the process of estimating the motion of the camera by examining the changes that the motion induces on the images made by it. This work consists of two parts. In the first part we propose a novel visual odometry algorithm. In the second part we describe a new data set.

The algorithm we propose exploits a scene structure typical for that seen by a moving car and is suitable for use in either the stereo or the monocular setting. We recover the rotation and the translation separately, thus dealing with two separate, smaller problems. The rotation is estimated by means of the infinite homography. The rotation estimation algorithm operates on distant image points using the 3-D to partition them into the distant and the near-by ones. We start with an initial estimate and then refine it using an iterative procedure. After the rotation is compensated for, the translation is found by means of the 1-point algorithm in the stereo setting and epipole computation for pure translational motion in the monocular setting.

We also created a new dataset that contains stereo video tracks synchronized with the DGPS locations of the vehicle. The dataset is suitable for future visual odometry experiments and is interesting because it has some challenging sequences (e.g. significant scene occlusions by a moving vehicle, fast camera motion, rural scenery, dusky sequences).

0.2 Summary of Our Work

In this work we present a novel algorithm for camera motion estimation. The novelty of the algorithm is in camera rotation estimation procedure. We rely on the fact that for scene points that are infinitely far from the camera, the motion of the projected (image) points may be described by an homography (the infinite homography). For distant points this assumption is nearly true. Our algorithm starts by partitioning the scene points into two sets: distant and near-by. Then, camera rotation is estimated from the distant points and, subsequently, the translation is recovered from the near-by points.

With respect to the classification of the visual odometry methods given in the introduction, our work is local, feature based, stereo odometry. We do not use bundle adjustment, however the results of our algorithm may be subsequently improved with some form of bundle adjustment.

The outline of the our method:

1. Feature detection. We use Harris corners.
2. Feature matching. The matching is done both across the stereo pair images as well as previous vs. current pair. We enforce epipolar constraint, chierality and use circle heuristics similar to reject outliers.
3. Partition the scene points into two sets: distant and near-by.

4. Estimate the rotation of the camera from the distant points.
5. Estimate the translation of the camera from the near-by points.

We choose the StereoScan as our baseline (our implementation of their work). The results in the show that on the KITTI dataset our rotation estimation method outperforms the baseline.

We also build a dataset that contains a number of sequences that may be used to benchmark visual odometry algorithms. We install synchronized stereo pair along a (D)GPS receiver on a car that travels in an urban as well as rural areas. Our goal is to cover some challenging conditions that happen in day-to-day driving, e.g., field of view occlusions, poor lighting conditions and low texture areas.

Note, that budget constraints led us to a simplified system that is capable of producing 3DOF data (i.e., location only) being unable to track vehicle orientation.